1   Definitions and disambiguations

In this section, we recall or rigorously redefine notions related to tensors, neural networks and graphs. Vector spaces considered below are assumed to be finite-dimensional and over the field of real numbers $\mathbb{R}$.

1.1   Tensors

Intuitively, tensors in the field of deep learning are defined as a generalization of vectors and matrices, as if vectors were tensors of rank 1 and matrices were tensors of rank 2. That is, they are objects in a linear space and their dimensions are indexed using as many indices as their rank, so that they can be represented by multidimensional arrays. In mathematics, a tensor is rather defined as a special class of multilinear functions. As such, a mathematical tensor is entirely defined on the cartesian product of the canonical bases onto which its outputs can be represented by a multidimensional array. In that sense, both definitions rejoin on their representation, but the underlying objects are different.
Our definition of tensors is such that they are a bit more than multidimensional arrays but not as much as mathematical tensors, for that they are embedded in a linear space and any deep learning object can be later define rigorously.

Definition 1.1. Tensor space
We define a *tensor space* $\mathbb{T}$ of rank r as a cartesian product of r vector spaces under the coordinate-wise sum and the mono-linear outer product.
Its *shape* is denoted $n_1 \times n_2 \times \cdots \times n_r$, where the $\{n_k\}$ are the dimensions of the vector spaces.

Hence, a tensor space is a linear space. Note that for naming conveniency, we distinguish between the terms *linear space* and *vector space*. That is, we abusively use the term *vector space* only to refer to a linear space that can be seen as a tensor space of rank 1, and *matrix space* to refer to a tensor space of rank 2. If there is no notion of rank defined, we rather use the term *linear space*. We also make a clear distinction between the term *dimension* (that is, for a tensor space it is equal to $\prod_{k=1}^{r} n_k$) and the term *rank* (equal to r).

Definition 1.2. Tensor
A *tensor* t is an object of a tensor space. The *shape* of t, which is the same as the shape of the tensor space it belongs to, is denoted $n_1^{(t)} \times n_2^{(t)} \times \cdots \times n_r^{(t)}$.

Definition 1.3. Indexing
An *entry* of a tensor t is a scalar denoted $t[i_1, i_2, \ldots, i_r]$.
A *subtensor* t' is a tensor of same rank composed of entries of t that are contiguous in the indexing, with at least one entry per rank. We denote $t' = t[l_1{:}u_1, l_2{:}u_2, \ldots, l_r{:}u_r]$, where the $\{l_p\}$ and the $\{u_p\}$ are the lower and upper bounds of the indices used by the entries that compose t'.

When using an index i_k for an entry of a tensor t, we implicitly assume that $i_k \in \{1, 2, \cdots, n_k^{(t)}\}$ if nothing is precised. For subtensor indexings, we don't necessarily write the lower bound index if it is equal to 1, neither the upper bound index if it is equal to $n_p^{(t)}$.

Definition 1.4. Slicing

A *slice* operation, along the last ranks $\{r_1, r_2, \ldots, r_s\}$, and indexed by $(i_{r_1}, i_{r_2}, \ldots, i_{r_s})$, is a morphism $s : \mathbb{T} = \prod_{k=1}^{r} \mathbb{V}_k \to \prod_{k=1}^{r-s} \mathbb{V}_k$, such that:

$$s(t)[i'_1, i'_2, \ldots, i'_{r-s}] = t[i'_1, i'_2, \ldots, i'_{r-s}, i_{r_1}, i_{r_2}, \ldots, i_{r_s}]$$
$$\text{i.e.} \quad s(t) := t[:, :, \ldots, :, i_{r_1}, i_{r_2}, \ldots, i_{r_s}]$$

where $:=$ means that values of the right operand are assigned to the left operand. We denote $t_{i_{r_1}, i_{r_2}, \ldots i_{r_s}}$ and call it the *slice* of t. Slicing along a subset of ranks that are not the lasts is defined similarly. $s(\mathbb{T})$ is called a *sliced subspace*.

Definition 1.5. Flattening

A *flatten* operation is an isomorphism $f : \mathbb{T} \to \mathbb{V}$, between a tensor space $\mathbb{T}$ of rank r and an n-dimensional vector space $\mathbb{V}$, where $n = \prod_{k=1}^{r} n_k$. It is characterized by a bijection in the index spaces $g : \prod_{k=1}^{r} \{1, \ldots, n_k\} \to \{1, \ldots, n\}$ such that

$$\forall t \in \mathbb{T}, f(t)[g(i_1, i_2, \ldots, i_r)] = f(t[i_1, i_2, \ldots, i_r])$$

An inverse operation is called a *de-flatten* operation.

Remark 1. **Row major ordering**

The choice of g determines in which order the indexing is made. g is reminescent of how data of multidimensional arrays or tensors are stored internally by programming languages. In most tensor manipulation languages, incrementing the memory adress (i.e. the output of g) will increment only i_r if $i_r < n_r$ (and then ranks are ordered in reverse lexicographic order to decide what index is also incremented). This is called *row major ordering*, as opposed to *column major ordering*. That is, in row major, g is defined as

$$g(i_1, i_2, \ldots, i_r) = \sum_{p=1}^{r} \left(\prod_{k=p+1}^{r} n_k \right) i_p \tag{1}$$

Definition 1.6. Reshaping

A *reshape* operation is an isomorphism defined on a tensor space $\mathbb{T} = \prod_{k=1}^{r} \mathbb{V}_k$ such that some of its basis vector spaces $\{\mathbb{V}_k\}$ are de-flattened and some of its sliced subspaces are flattened.

2

Definition 1.7. Contraction
A *tensor contraction* between two tensors, along ranks of same dimensions, is defined by natural extension of the dot product operation to tensors.
More precisely, let $\mathbb{T}_1$ a tensor space of shape $n_1^{(1)} \times n_2^{(1)} \times \cdots \times n_{r_1}^{(1)}$, and $\mathbb{T}_2$ a tensor space of shape $n_1^{(2)} \times n_2^{(2)} \times \cdots \times n_{r_2}^{(2)}$, such that $\forall k \in \{1, 2, \ldots, s\}, n_{r_1-(s-k)}^{(1)} = n_k^{(2)}$, then the tensor contraction between $t_1 \in \mathbb{T}_1$ and $t_2 \in \mathbb{T}_2$ is defined as:

$$\begin{cases} t_1 \otimes t_2 = t_3 \in \mathbb{T}_3 \text{ of shape } n_1^{(1)} \times \cdots \times n_{r_1-s}^{(1)} \times n_{s+1}^{(2)} \times \cdots \times n_{r_2}^{(2)} \text{ where} \\ t_3[i_1^{(1)}, \ldots, i_{r_1-s}^{(1)}, i_{s+1}^{(2)}, \ldots, i_{r_2}^{(2)}] = \\ \sum_{k_1=1}^{n_1^{(2)}} \cdots \sum_{k_s=1}^{n_s^{(2)}} t_1[i_1^{(1)}, \ldots, i_{r_1-s}^{(1)}, k_1, \ldots, k_s] \, t_2[k_1, \ldots, k_s, i_{s+1}^{(2)}, \ldots, i_{r_2}^{(2)}] \end{cases}$$

For the sake of simplicity, we omit the case where the contracted ranks are not the last ones for t_1 and the first ones for t_2. But this definition still holds in the general case subject to a permutation of the indices.

Remark 2. **Covariant and contravariant indices**
Indices of the left handed operand t_1 that are not contracted are called *covariant* indices. Those that are contracted are called *contravariant* indices. For the right operand t_2, the naming convention is the opposite. Using subscript notation for covariant indices and superscript notation for contravariant indices, the previous tensor contraction can be written using the Einstein summation convention as:

$$t_1{}_{i_1^{(1)}\cdots i_{r_1-s}^{(1)}}{}^{k_1 \cdots k_s} \otimes t_2{}_{k_1 \cdots k_s}{}^{i_{s+1}^{(2)}\cdots i_{r_2}^{(2)}} = t_3{}_{i_1^{(1)}\cdots i_{r_1-s}^{(1)}}{}^{i_{s+1}^{(2)}\cdots i_{r_2}^{(2)}} \tag{2}$$

Dot product $u_k v^k = \lambda$ and matrix product $A_i{}^k B_k{}^j = C_i{}^j$ are common examples of tensor contractions.

Remark 3. **Matrix product equivalence**
Using reshapings that groups all covariant indices into a single index and all contravariant indices into another single index, any tensor contraction can be rewritten as a matrix product.

Proof. Using notation of (2), with the reshapings $t_1 \mapsto T_1$, $t_2 \mapsto T_2$ and $t_3 \mapsto T_3$ we can rewrite

$$T_1{}_{g_i(i_1^{(1)}, \ldots, i_{r_1-s}^{(1)})}{}^{g_k(k_1, \ldots, k_s)} \otimes T_2{}_{g_k(k_1, \ldots, k_s)}{}^{g_j(i_{s+1}^{(2)}, \ldots, i_{r_2}^{(2)})} = T_3{}_{g_i(i_1^{(1)}, \ldots, i_{r_1-s}^{(1)})}{}^{g_j(i_{s+1}^{(2)}, \ldots, i_{r_2}^{(2)})}$$

where g_i, g_k and g_j are bijections defined similarly as in (1). $\qquad \square$

Definition 1.8. Convolution
The *n-dimensional convolution*, denoted $*_n$, between $t_1 \in \mathbb{T}_1$ and $t_2 \in \mathbb{T}_2$, where $\mathbb{T}_1$ and $\mathbb{T}_2$ are of the same rank n such that $\forall p \in \{1, 2, \ldots, n\}, n_p^{(1)} \geq n_p^{(2)}$, is

defined as:

$$
\begin{cases}
t_1 *_n t_2 = t_3 \in \mathbb{T}_3 \text{ of shape } n_1^{(3)} \times \cdots \times n_n^{(3)} \text{ where} \\
\forall p \in \{1, 2, \ldots, n\}, n_p^{(3)} = n_p^{(1)} - n_p^{(2)} + 1 \\
t_3[i_1, \ldots, i_n] = \sum_{k_1=1}^{n_1^{(3)}} \cdots \sum_{k_n=1}^{n_n^{(3)}} t_1[i_1 + k_1 - 1, \ldots, i_n + k_n - 1]\, t_2[k_1, \ldots, k_n]
\end{cases}
$$

Definition 1.9. Pooling

TODO: Use indexing, maybe add stride in this subsection rather than in the next

1.2   Neural Networks

We denote by I_f the *domain of definition* of a function f ("I" for "input") and by $O_f = f(I_f)$ its *image* ("O" for "output"), and we represent it as $I_f \xrightarrow{f} O_f$. An *activation function* h defined from a tensor space to itself is defined as an univariate function applied dimension-wise and we use the functional notation $h(v)[i_1, i_2, \ldots, i_r] = h(v[i_1, i_2, \ldots, i_r])$.

Definition 1.10. Neural network
 Let F be a function such that I_f and O_f are vector or tensor spaces.
F is a *functional formulation* of a *neural network* if there are a series of linear or affine functions $(g_k)_{k=1,2,..,L}$ and a series of non-linear derivable activation functions $(h_k)_{k=1,2,..,L}$ such that:

$$
\begin{cases}
\forall k \in \{1, 2, .., L\}, f_k = h_k \circ g_k, \\
I_F = I_{f_1} \xrightarrow{f_1} O_{f_1} \cong I_{f_2} \xrightarrow{f_2} \ldots \xrightarrow{f_L} O_{f_L} = O_F, \\
F = f_L \circ \ldots \circ f_2 \circ f_1
\end{cases}
$$

The couple (g_k, h_k) is called the *k-th layer* of the neural network. For $x \in I_f$, we denote by $x_k = f_k \circ \ldots \circ f_2 \circ f_1(x)$ the *activations* of the k-th layer.

Remark 4. **Bias**
Affine functions $\widetilde{g}$ can be written as a sum between a linear function g and a constant vector b. For notational conveniency, we only consider linear functions in this section. b is called the *bias*. Its role is to augment the expressivity of the neural network's family of functions.

Definition 1.11. Connectivity matrix
Let g a linear function. Without loss of generality subject to a flattening, let's suppose I_g and O_g are vector spaces. Then there exists a *connectivity matrix* W_g, such that:

$$
\forall x \in I_g, g(x) = W_g x
$$

We denote W_k the connectivity matrix of the k-th layer.

Remark 5. **Biological inspiration**
A *(computational) neuron* is a computational unit that is biologically inspired. Each neuron should be capable of:

1. receiving modulated signals from other neurons and aggregate them,

2. applying to the result a derivable activation,

3. passing the signal to other neurons.

That is to say, each domain $\{I_{f_k}\}$ and O_F can be interpreted as a layer of neurons, with one neuron for each dimension. The connectivity matrices $\{W_k\}$ describe the connexions between each successive layers. A neuron is illustrated on Figure 1.


Figure 1: A neuron

Definition 1.12. Weights
Let consider the k-th layer of a neural networks. We define its weights as coordinates of a vector θ_k, called the *weight kernel*, such that:

$$\forall (i,j), \begin{cases} \exists p, W_k[i,j] := \theta_k[p] \\ \text{or } W_k[i,j] = 0 \end{cases}$$

A weight p that appears multiple times in W_k is said to be *shared*. Two parameters of W_k that share a same weight p are said to be *tied*. The number of weights of the k-th layer is $n_1^{(\theta_k)}$.

Remark 6. **Training**
A *loss* function $\mathcal{L}$ penalizes the output $x_L = F(x)$ relatively to what can be expected. Gradient w.r.t. θ_k, denoted $\vec{\nabla}_{\theta_k}$, is used to update the weights via an optimization algorithm based on gradient descent and a learning rate α, that is:

$$\theta_k^{(\text{new})} = \theta_k^{(\text{old})} - \alpha \cdot \vec{\nabla}_{\theta_k} \left(\mathcal{L}\left(x_L, \theta_k^{(\text{old})}\right) + R\left(\theta_k^{(\text{old})}\right) \right) \tag{3}$$

where α can be a scalar or a vector, $\cdot$ can denote outer or pointwise product, and R is a regularizer. They depend on the optimization algorithm.

Remark 7. **Linear complexity**
The complexity of computing the gradients is linear with the number of weights.

Proof. Thanks to the chain rule, $\vec{\nabla}_{\theta_k}$ can be computed using gradients that are w.r.t. x_k, denoted $\vec{\nabla}_{x_k}$, which in turn can be computed using gradients w.r.t. outputs of the next layer $k+1$, up to the gradients given on the output layer. That is:

$$\vec{\nabla}_{\theta_k} = J_{\theta_k}(x_k)\vec{\nabla}_{x_k} \tag{4}$$

$$\vec{\nabla}_{x_k} = J_{x_k}(x_{k+1})\vec{\nabla}_{x_{k+1}}$$

$$\vec{\nabla}_{x_{k+1}} = J_{x_{k+1}}(x_{k+2})\vec{\nabla}_{x_{k+2}} \tag{5}$$

$$\dots$$

$$\vec{\nabla}_{x_{L-1}} = J_{x_{L-1}}(x_L)\vec{\nabla}_{x_L}$$

Obtaining,

$$\vec{\nabla}_{\theta_k} = J_{\theta_k}(x_k)(\prod_{p=k}^{L-1} J_{x_p}(x_{p+1}))\vec{\nabla}_{x_L} \tag{6}$$

where $J_{\mathrm{wrt}}(.)$ are the respective jacobians which can be determined with the layer's expressions and the $\{x_k\}$; and $\vec{\nabla}_{x_L}$ can be determined using $\mathcal{L}$, R and x_L. $\qquad\square$

This allows to compute the gradients with a complexity that is linear with the number of weights (only one computation of the $\{x_k\}$), instead of being quadratic if it were done with the difference quotient expression of the derivatives (one more computation of the $\{x_k\}$ per weight).

Remark 8. **Back propagation**
We can remark that (5) rewrites as

$$\vec{\nabla}_{x_k} = J_{x_k}(x_{k+1})\vec{\nabla}_{x_{k+1}}$$

$$= J_{x'_k}(h(x'_k))J_{x_k}(W_k x_k)\vec{\nabla}_{x_{k+1}} \tag{7}$$

where $x'_k = W_k x_k$, and these jacobians can be expressed as:

$$J_{x'_k}(h(x'_k))[i, j] = \delta_i^j h'(x'_k[i])$$

$$J_{x'_k}(h(x'_k)) = I\, h'(x'_k) \tag{8}$$

$$J_{x_k}(W_k x_k) = W_k^T \tag{9}$$

That means that we can write $\vec{\nabla}_{x_k} = (\widetilde{h}_k \circ \widetilde{g}_k)(\vec{\nabla}_{x_{k+1}})$ such that the connectivity matrix $\widetilde{W}_k$ is obtained by transposition. This can be interpreted as gradient calculation being a *back-propagation* on the same neural network, in opposition of the *forward-propagation* done to compute the output.

Definition 1.13. Connections
The set of *connections* of a layer (g, h), denoted C_g, is defined as:

$$C_g = \{(i,j), \exists p, W_g[i,j] := \theta_g[p]\}$$

We have $0 \le |C_g| \le n_1^{(W_g)} n_2^{(W_g)}$.

Definition 1.14. Dense layer
A *dense layer* (g, h) is a layer such that $|C_g| = n_1^{(W_g)} n_2^{(W_g)}$, i.e. all possible connections exist. The map $(i,j) \mapsto p$ is usually a bijection, meaning that there is no weight sharing.

Definition 1.15. Partially connected layer
A *partially connected layer* (g, h) is a layer such that $|C_g| < n_1^{(W_g)} n_2^{(W_g)}$.
A *sparsely connected layer* (g, h) is a layer such that $|C_g| \ll n_1^{(W_g)} n_2^{(W_g)}$.

Definition 1.16. Convolutional layer
A *n-dimensional convolutional layer* (g, h) is such that the weight kernel θ_g can be reshaped into a tensor w of rank $n + 2$, and such that

$$\begin{cases} I_g \text{ and } O_g \text{ are tensor spaces of rank } n+1 \\ \forall x \in I_g, g(x) = (g(x)_q = \sum_p x_p *_n w_{p,q})_{\forall q} \end{cases}$$

where p and q index slices along the last ranks. The slices $g(x)_q$ are typically called *feature maps*.

Remark 9. **Connectivity matrix of a convolution**
Note that a convolutional layer (g, h) is equivalently defined as W_g being a $n_{n+1}^{(I_g)} \times n_{n+1}^{(O_g)}$ block matrix such that its blocks are Toeplitz matrices.

Proof. Let's recall from definition 1.8 that

$$y_q = x_p *_n w_{p,q}[i_1, \ldots, i_n]$$

$$= \sum_{k_1=1}^{n_1^{(x_p)} - n_1^{(w_{p,q})}+1} \cdots \sum_{k_n=1}^{n_n^{(x_p)} - n_n^{(w_{p,q})}+1} x_p[i_1 + k_1 - 1, \ldots, i_n + k_n - 1] \, w_{p,q}[k_1, \ldots, k_n]$$

$$= \sum_{j_1=i_1}^{i_1 + n_1^{(x_p)} - n_1^{(w_{p,q})}} \cdots \sum_{j_n=i_n}^{i_n + n_n^{(x_p)} - n_n^{(w_{p,q})}} x_p[j_1, \ldots, j_n] \, w_{p,q}[j_1 - i_1 + 1, \ldots, j_n - i_n + 1]$$

where $\forall p, j_p = i_p + k_p - 1$

Let a reshaping $w_{p,q} \mapsto \widetilde{w}_{p,q}$ such that its ranks are de-flattened using the inverses of the bijections $\{g_t : (i_t, j_t) \mapsto j_t - i_t + 1\}$. Then, subject to an index permutation and using notations as in (2), we can rewrite:

$$y_{q_{i_1 \cdots i_n}} = x_{p_{j_1 \cdots j_n}} \widetilde{w}_{p,q} {}^{j_1 \cdots j_n}{}_{i_1 \cdots i_n}$$

where we recognize that the n-d convolution is expressed as a tensor contraction. Using the matrix product equivalence from remark 3, we can reshape $\widetilde{w}_{p,q}$ to a matrix $W_{p,q}$ and obtain the block (p,q) of W_g.

The Toeplitz property holds for $\widetilde{w}_{p,q}$ for couples of covariant and contravariant indices (i_t, j_t) because incrementing them both keeps the weight index k_t constant. Hence, $W_{p,q}$ is a Toeplitz matrix. The proof of the converse also holds because reshapings are isomorphisms. □

Definition 1.17. Convolutional layer with stride
Let g the linear part of a convolution layer with *stride* $s_p > 1$ along the p-th rank, and $\widetilde{g}$ the same linear part if it was a regular convolutional layer as defined above. Then g is defined as

$$\forall i_p \in \{1, \lfloor \frac{n_p^{(g(x))}}{s_p} \rfloor\}, \forall x \in I_g, g(x)_{i_p} = \widetilde{g}(x)_{s_p i_p}$$

where i_p and $s_p i_p$ index slices along the p-th rank.

TODO: below

Definition 1.18. Pooling
A layer with *pooling* (g, h) is such that $g = g_1 \circ g_2$, where (g_1, h) is a layer and g_2 is a pooling operation.

A layer with *dropout* (g, h) is such that $h = h_1 \circ h_2$, where (g, h_2) is a layer and h_1 is a dropout operation [?]. When dropout is used, a certain number of neurons are randomly set to zero during the training phase, compensated at test time by scaling down the whole layer. This is done to prevent overfitting.
TODO: rephrase

A multilayer perceptron (MLP) [?] is a neural network composed of only dense layers. A convolutional neural network (CNN) [?] is a neural network composed of convolutional layers.

Neural networks are commonly used for machine learning tasks. For example, to perform supervised classification, we usually add a dense output layer $s = (g_{L+1}, h_{L+1})$ with as many neurons as classes. We measure the error between an output and its expected output with a discriminative loss function $\mathcal{L}$. During the training phase, the weights of the network are adapted for the classification task based on the errors that are back-propagated [?] via the chain rule and according to a chosen optimization algorithm (e.g. [?]).

1.3 Graphs

TODO: check this subsection

A graph G is defined as a couple (V, E) where V represents the set of nodes and $E \subseteq \binom{V}{2}$ is the set of edges connecting these nodes.

We encounter the notion of graphs several times in deep learning:

- Connections between two layers of a deep learning model can be represented as a bipartite graph, coined *connectivity graph*. It encodes how the information is propagated through a layer to another. See section 1.3.1.

- A computation graph is used by deep learning frameworks to keep track of the dependencies between layers of a deep learning models, in order to compute forward and back-propagation. See section 1.3.2.

- A graph can represent the underlying structure of an object (often a vector), whose nodes represent its features. See section 1.3.3.

- Datasets can also be graph-structured, where the nodes represent the objects of the dataset. See section 1.3.4.

1.3.1 Connectivity graph

A Connectivity graph is the bipartite graph whose adjacency matrix is the connectivity matrix of a layer of neurons. Formally, given a linear part of a layer, let $\mathbf{x}$ and $\mathbf{y}$ be the input and output signals, n the size of the set of input neurons $N = \{u_1, u_2, \ldots, u_n\}$, and m the size of the set of output neurons $M = \{v_1, v_2, \ldots, v_m\}$. This layer implements the equation $y = \Theta x$ where Θ is a $n \times m$ matrix.

Definition 1.19. The *connectivity graph* $G = (V, E)$ is defined such that $V = N \cup M$ and $E = \{(u_i, v_j) \in N \times M, \Theta_{ij} \neq 0\}$.

I.e. the connectivity graph is obtained by drawing an edge between neurons for which $\Theta_{ij} \neq 0$. For instance, in the special case of a complete bipartite graph, we would obtain a dense layer. Connectivity graphs are especially useful to represent partially connected layers, for which most of the Θ_{ij} are 0. For example, in the case of layers characterized by a small local receptive field, the connectivity graph would be sparse, and output neurons would be connected to a set of input neurons that corresponds to features that are close together in the input space. Figure 2 depicts some examples.

Connectivity graphs also allow to graphically modelize how weights are tied in a neural layer. Let's suppose the Θ_{ij} are taking their values only into the finite set $K = \{w_1, w_2, \ldots, w_\kappa\}$ of size κ, which we will refer to as the *kernel* of *weights*. Then we can define a labelling of the edges $s : E \rightarrow K$. s is called the *weight sharing scheme* of the layer. This layer can then be formulated as $\forall v \in M, y_v = \sum_{u \in N, (u,v) \in E} w_{s(u,v)} x_u$. Figure 3 depicts the connectivity graph of a 1-d convolution layer and its weight sharing scheme.

Figure 2: Examples

Figure 3: Depiction of a 1D-convolutional layer and its weight sharing scheme.

TODO: Add weight sharing scheme in Figure 3

1.3.2 Computation graph

1.3.3 Underlying graph structure

1.3.4 Graph-structured dataset

transductive vs inductive

1.4 Geometric grids

1.5 Grid graphs

1.6 Spatial graphs

1.7 Projections of spatial graphs